

# Policy Iteration

## Reinforcement Learning

Roberto Capobianco



SAPIENZA  
UNIVERSITÀ DI ROMA

# Recap

# Bellman Equation

---

The value of a certain state is expanded in terms of the current reward and the value of the next states according to the policy

$r$  here is function of  $s$  and  $\pi(s)$

$$V^\pi(s_t) = \mathbb{E}_\pi[r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \gamma^3 r_{t+3} + \dots | s_t] = r_t + \gamma \mathbb{E}_{s' \sim p(\cdot | s, \pi(s))} [V^\pi(s')]$$

$$Q^\pi(s_t, a) = r_t + \gamma \mathbb{E}_{s' \sim p(\cdot | s, a)} [V^\pi(s')]$$

$r$  here is function of  $s$  and  $a$

As a result  $V(s) = Q(s, \pi(s))$




# Bellman Optimality Example

---

$$V^*(s) = \max_a [r(s, a) + \gamma \mathbb{E}_{s' \sim p(\cdot | s, a)} V^*(s')] ]$$

- Try  $a'$ , get  $r(s, a')$ ,  
compute  
 $Q^*(s, a') = r(s, a') + \gamma V^*(s')$
- Try  $a''$ , get  $r(s, a'')$ ,  
compute  
 $Q^*(s, a'') = r(s, a'') + \gamma V^*(s'')$



Assume we know  $V^*$  at  
 $s'$  and  $s''$

$$V^*(s) = \max_{a', a'', \dots} \{ Q^*(s, a'), Q^*(s, a'') \}$$



# Exact Policy Evaluation

---  
We know that **for ALL states**, Bellman equation holds

$$V^\pi(s) = r + \gamma \mathbb{E}_{s' \sim p(\cdot | s, \pi(s))} [V^\pi(s')]$$

We can combine all the constraints together:

Since we have this set of constraints

$$V = R + \gamma P V$$

we can solve for  $V$  as

$$V = (I - \gamma P)^{-1} R$$

The diagram shows the Bellman equation in matrix notation. On the left is a vertical vector labeled  $V$  with a red label  $V(s)$  in its middle cell. This is followed by an equals sign, then a vertical vector labeled  $R$  with a red label  $r(s, \pi(s))$  in its middle cell. This is followed by a plus sign, then a scalar  $\gamma$  (in black), then a matrix labeled  $P$  with a red label  $P(\cdot | s, \pi(s))$  in its middle row, and finally a vertical vector labeled  $V$ .

:( Nice but computationally expensive: inverting the matrix is  $O(S^3)$



# Fixed-Point Iteration & Contractions

— — —

What is a fixed-point? A point where holds

$$x = f(x)$$

How can we find such points?

- Initialize  $x_0$
- Repeat  $x_{i+1} = f(x_i)$
- Stop at convergence where  $x$  is found and does not change anymore

Convergence to a fixed-point is possible thanks to the existence of **contraction mappings**

$f: M \rightarrow M$  ( $M$  is a metric space) is a contraction mapping if:

$$|f(x) - f(x')| \leq k|x - x'| \text{ for } k \text{ in } [0, 1)$$



# Iterative Policy Evaluation

---

- Initialize  $V_0$  in  $[0, 1/(1-\gamma)]$  (typically 0)
- Until convergence:

$$V_{i+1} = R + \gamma P V_i$$

(note: this is using matricial form because it's doing it for all states)

$$\|V^{t+1} - V^\pi\|_\infty \leq \gamma \|V^t - V^\pi\|_\infty \leq \gamma^{t+1} \|V^0 - V^\pi\|_\infty$$

For each iteration it's  $O(S^2)$



# How to Find the Optimal Policy?

— — —

Now, what we're really interested in is finding the optimal policy  $\pi^*$

**Let's use Bellman optimality!** ...and the Bellman Operator (which is a contraction)

$$TQ(s,a) = r(s,a) + \gamma E_{s' \sim p(\cdot | s, a)} \max_{a'} [Q(s', a')]$$

Since  $Q: S \times A \rightarrow \mathbb{R}$ , then also  $TQ: S \times A \rightarrow \mathbb{R}$





# Value Iteration & Optimal Policy

---

We can obtain  $Q^* = TQ^*$ , since  $Q^*$  is a fixed-point solution to  $Q = TQ$

- Initialize  $\|Q_0\|$  in  $[0, 1/(1-\gamma)]$  (typically 0)
- Until convergence, for all states and actions:

$$Q_{i+1} = TQ_i$$

$$\|Q_{i+1} - Q^*\| = \|TQ_i - TQ^*\| \leq \gamma \|Q_i - Q^*\| \leq \gamma^{i+1} \|Q_0 - Q^*\|$$

We know that  $\pi^*(s) = \operatorname{argmax}_a Q^*(s,a)$ , and since  $Q_i(s,a) \approx Q^*(s,a)$  we could choose

$$\pi_i(s) = \operatorname{argmax}_a Q_i(s,a)$$



# End - Recap



SAPIENZA  
UNIVERSITÀ DI ROMA

# Clarification: $V_0$ Initialization

---

We Initialize  $V_0$  in  $[0, 1/(1-\gamma)]$  (typically 0)

**WHY?**



# Clarification: $V_0$ Initialization

---

We Initialize  $V_0$  in  $[0, 1/(1-\gamma)]$  (typically 0)

**WHY?**

**Under the assumption that  $R(s,a)$  is in  $[0, 1]$**

The maximum possible value of  $V$  is  $1 + \gamma + \gamma^2 + \dots = 1/(1-\gamma)$



# Clarification: $V_0$ Initialization

---

We Initialize  $V_0$  in  $[0, 1/(1-\gamma)]$  (typically 0)

**WHY?**

**Under the assumption that  $R(s,a)$  is in  $[0, 1]$**

The maximum possible value of  $V$  is  $1 + \gamma + \gamma^2 + \dots = 1/(1-\gamma)$

There is no reason to initialize higher than that!



# Clarification: Number of Iterations of Value Iteration

---

If we want an  $\epsilon$  error on the quality of the policy, to determine the number of iterations  $i$  we can just solve for it in this equation

$$2\gamma^i / (1-\gamma) \|Q_0 - Q^*\| \leq \epsilon$$



# Clarification: Number of Iterations of Value Iteration

---

If we want an  $\epsilon$  error on the quality of the policy, to determine the number of iterations  $i$  we can just solve for it in this equation

$$2\gamma^i / (1-\gamma) \|Q_0 - Q^*\| \leq \epsilon$$

In range  $[0, 1/(1-\gamma)]$



# Clarification: Number of Iterations of Value Iteration

---

If we want an  $\epsilon$  error on the quality of the policy, to determine the number of iterations  $i$  we can just solve for it in this equation

$$2\gamma^i / (1-\gamma) \|Q_0 - Q^*\| \leq \epsilon$$

In range  $[0, 1/(1-\gamma)]$





# Clarification: Number of Iterations of Value Iteration

---

If we want an  $\epsilon$  error on the quality of the policy, to determine the number of iterations  $i$  we can just solve for it in this equation

$$2\gamma^i / (1-\gamma) \|Q_0 - Q^*\| \leq \epsilon$$

By the infinity norm, the maximum value is  $1/(1-\gamma)$



# Clarification: Number of Iterations of Value Iteration

---

If we want an  $\epsilon$  error on the quality of the policy, to determine the number of iterations  $i$  we can just solve for it in this equation

$$2\gamma^i / (1-\gamma) \|Q_0 - Q^*\| = 2(1 - (1-\gamma)^i) / (1-\gamma) \|Q^*\| \leq \epsilon$$

Additionally add and  
subtract 1



# Clarification: Number of Iterations of Value Iteration

---

If we want an  $\epsilon$  error on the quality of the policy, to determine the number of iterations  $i$  we can just solve for it in this equation

$$2\gamma^i / (1-\gamma) \|Q_0 - Q^*\| = 2(1 - (1-\gamma))^i / (1-\gamma) \|Q^*\| \leq \epsilon$$

$$2(1 - (1-\gamma))^i / (1-\gamma) \|Q^*\| \leq 2e^{-(1-\gamma)i} / (1-\gamma)^2 \leq \epsilon$$

$1+x \leq e^x$  for all reals



# Clarification: Number of Iterations of Value Iteration

---

If we want an  $\epsilon$  error on the quality of the policy, to determine the number of iterations  $i$  we can just solve for it in this equation

$$2\gamma^i / (1-\gamma) \|Q_0 - Q^*\| = 2(1-(1-\gamma))^i / (1-\gamma) \|Q^*\| \leq \epsilon$$

$$2(1-(1-\gamma))^i / (1-\gamma) \|Q^*\| \leq 2e^{-(1-\gamma)i} / (1-\gamma)^2 \leq \epsilon$$

$$e^{-(1-\gamma)i} \leq \epsilon(1-\gamma)^2 / 2$$

$$\log(1/x) = -\log(x)$$

$$-i(1-\gamma) \leq -\log(2 / (\epsilon(1-\gamma)^2))$$



# Clarification: Number of Iterations of Value Iteration

---

If we want an  $\epsilon$  error on the quality of the policy, to determine the number of iterations  $i$  we can just solve for it in this equation

$$2\gamma^i / (1-\gamma) \|Q_0 - Q^*\| = 2(1 - (1-\gamma)^i) / (1-\gamma) \|Q^*\| \leq \epsilon$$

$$2(1 - (1-\gamma)^i) / (1-\gamma) \|Q^*\| \leq 2e^{-(1-\gamma)^i} / (1-\gamma)^2 \leq \epsilon$$

$$e^{-(1-\gamma)^i} \leq \epsilon(1-\gamma)^2 / 2$$

$$-i(1-\gamma) \leq -\log(2 / (\epsilon(1-\gamma)^2))$$

$$i \geq \log(2 / (\epsilon(1-\gamma)^2)) / (1-\gamma)$$



# Clarification: Number of Iterations of Value Iteration

---

If we want an  $\epsilon$  error on the quality of the policy, to determine the number of iterations  $i$  we can just solve for it in this equation

$$i \geq \frac{\log \frac{2}{\epsilon(1-\gamma)^2}}{1-\gamma}$$



# Another Note on Value Iteration

— — —

- $Q_t$  is approximating  $Q^*$
- From  $Q_t$  we compute a policy  $\pi_t$

However...



# Another Note on Value Iteration

---

- $Q_t$  is approximating  $Q^*$
- From  $Q_t$  we compute a policy  $\pi_t$

However...

**$Q_t$  is generally different from  $Q^{\pi_t}$  until we converge to approximately  $Q^*$**

E.g,  $Q_0$  is just a random initial guess, maybe not corresponding to any policy's  $Q$  value





# Complexity of Value Iteration

— — —

For each iteration it's  $O(S^2A)$



# Policy Iteration

— — —

- Outputs policies at every iteration:  $\{\pi_0, \pi_1, \pi_2 \dots \pi_T\}$
- Different from Value Iteration that was outputting values



# Policy Iteration

---

- Outputs policies at every iteration:  $\{\pi_0, \pi_1, \pi_2 \dots \pi_T\}$
- Different from Value Iteration that was outputting values

Procedure:

1. Start with a random guess  $\pi_0$  (can be deterministic or stochastic)
2. For  $t=0, \dots, T$ :

$$Q^\pi(s_t, a) = r_t + \gamma \mathbb{E}_{s' \sim p(\cdot | s, a)} [V^\pi(s')]$$

  - a. Do **policy evaluation** and compute  $Q^{\pi^t}$  for all  $s, a$
  - b. Do **policy improvement** as  $\pi_{t+1} = \operatorname{argmax}_a Q^{\pi^t}(s, a)$  for all  $s$



# Policy Iteration

---

- Outputs policies at every iteration:  $\{\pi_0, \pi_1, \pi_2 \dots \pi_T\}$
- Different from Value Iteration that was outputting values

Procedure:

1. Start with a random guess  $\pi_0$  (can be deterministic or stochastic)



# Policy Iteration

- Outputs policies at every iteration:  $\{\pi_0, \pi_1, \pi_2 \dots \pi_T\}$
- Different from Value Iteration that was outputting values

Procedure:

1. Start with a random guess  $\pi_0$  (can be deterministic or stochastic)
2. For  $t=0, \dots, T$ :
  - a. Do **policy evaluation** and compute  $Q^{\pi^t}$  for all  $s, a$

Remember that  $Q^{\pi}(s_t, a) = r_t + \gamma \mathbb{E}_{s' \sim p(\cdot | s, a)} [V^{\pi}(s')]!$



# Policy Iteration

---

- Outputs policies at every iteration:  $\{\pi_0, \pi_1, \pi_2 \dots \pi_T\}$
- Different from Value Iteration that was outputting values

Procedure:

1. Start with a random guess  $\pi_0$  (can be deterministic or stochastic)
2. For  $t=0, \dots, T$ :
  - a. Do **policy evaluation** and compute  $Q^{\pi^t}$  for all  $s, a$

Remember that  $Q^{\pi}(s_t, a) = r_t + \gamma \mathbb{E}_{s' \sim p(\cdot | s, a)} [V^{\pi}(s')]$ !

We can first compute  $V$ , for example, and then get  $Q$  from that



# Policy Iteration

---

- Outputs policies at every iteration:  $\{\pi_0, \pi_1, \pi_2 \dots \pi_T\}$
- Different from Value Iteration that was outputting values

Procedure:

1. Start with a random guess  $\pi_0$  (can be deterministic or stochastic)
2. For  $t=0, \dots, T$ :
  - a. Do **policy evaluation** and compute  $Q^{\pi^t}$  for all  $s, a$

For simplicity and to forget about approximation errors, let's assume we use the exact policy iteration



# Policy Iteration

---

- Outputs policies at every iteration:  $\{\pi_0, \pi_1, \pi_2 \dots \pi_T\}$
- Different from Value Iteration that was outputting values

Procedure:

1. Start with a random guess  $\pi_0$  (can be deterministic or stochastic)
2. For  $t=0, \dots, T$ :
  - a. Do **policy evaluation** and compute  $Q^{\pi^t}$  for all  $s, a$

Differently from Value Iteration, here we are outputting  $Q$  values of actual policies!





# Policy Iteration

---

- Outputs policies at every iteration:  $\{\pi_0, \pi_1, \pi_2 \dots \pi_T\}$
- Different from Value Iteration that was outputting values

Procedure:

1. Start with a random guess  $\pi_0$  (can be deterministic or stochastic)
2. For  $t=0, \dots, T$ :
  - a. Do **policy evaluation** and compute  $Q^{\pi^t}$  for all  $s, a$
  - b. Do **policy improvement** as  $\pi_{t+1} = \arg\max_a Q^{\pi^t}(s, a)$  for all  $s$



# Policy Iteration

---

Procedure:

1. Start with a random guess  $\pi_0$  (can be deterministic or stochastic)
2. For  $t=0, \dots, T$ :
  - a. Do **policy evaluation** and compute  $Q^{\pi^t}$  for all  $s, a$
  - b. Do **policy improvement** as  $\pi_{t+1} = \operatorname{argmax}_a Q^{\pi^t}(s, a)$  for all  $s$

This algorithm only makes progress, and the performance progress of the policy is monotonic



# Properties of Policy Iteration


— — —

- Monotonic improvement:  $Q^{\pi^{t+1}} \geq Q^{\pi^t}$  for all  $s, a$
- Convergence:  $\|V^{\pi^i} - V^*\| \leq \gamma^{i+1} \|V^{\pi^0} - V^*\|$




# Properties of Policy Iteration

---



Credits: David Silver




# Monotonic Improvement

---

$$\pi_{t+1} = \operatorname{argmax}_a Q^{\pi_t}(s, a)$$

We want to show that  $Q^{\pi_{t+1}} \geq Q^{\pi_t}$  for all  $s, a$




# Monotonic Improvement

---

$$\pi_{t+1} = \operatorname{argmax}_a Q^{\pi_t}(s, a)$$

We want to show that  $Q^{\pi_{t+1}} \geq Q^{\pi_t}$  for all  $s, a$




We are back  
at the  
starting  
point: we can  
be recursive!



# Monotonic Improvement

---

We want to show that  $Q^{\pi_{t+1}} \geq Q^{\pi_t}$  for all  $s, a$



# Monotonic Improvement

$$\pi_{t+1} = \operatorname{argmax}_a Q^{\pi^t}(s, a)$$

---

We want to show that  $Q^{\pi^{t+1}} \geq Q^{\pi^t}$  for all  $s, a$

expand definition and simplify  $r(s, a)$

$$\begin{aligned} Q^{\pi^{t+1}}(s, a) - Q^{\pi^t}(s, a) &= \gamma \mathbb{E}_{s' \sim P(s, a)} \left[ Q^{\pi^{t+1}}(s', \pi^{t+1}(s')) - Q^{\pi^t}(s', \pi^t(s')) \right] \\ &= \gamma \mathbb{E}_{s' \sim P(s, a)} \left[ Q^{\pi^{t+1}}(s', \pi^{t+1}(s')) - Q^{\pi^t}(s', \pi^{t+1}(s')) + Q^{\pi^t}(s', \pi^{t+1}(s')) - Q^{\pi^t}(s', \pi^t(s')) \right] \\ &\geq \gamma \mathbb{E}_{s' \sim P(s, a)} \left[ Q^{\pi^{t+1}}(s', \pi^{t+1}(s')) - Q^{\pi^t}(s', \pi^{t+1}(s')) \right] \geq \dots, \geq -\gamma^\infty / (1 - \gamma) = 0 \end{aligned}$$





# Monotonic Improvement

$$\pi_{t+1} = \operatorname{argmax}_a Q^{\pi^t}(s, a)$$

---

We want to show that  $Q^{\pi^{t+1}} \geq Q^{\pi^t}$  for all  $s, a$

$$\begin{aligned} Q^{\pi^{t+1}}(s, a) - Q^{\pi^t}(s, a) &= \gamma \mathbb{E}_{s' \sim P(s, a)} \left[ Q^{\pi^{t+1}}(s', \pi^{t+1}(s')) - Q^{\pi^t}(s', \pi^t(s')) \right] \\ &\quad \text{add and subtract the Q of our intermediate policy} \\ &= \gamma \mathbb{E}_{s' \sim P(s, a)} \left[ Q^{\pi^{t+1}}(s', \pi^{t+1}(s')) - \underbrace{Q^{\pi^t}(s', \pi^{t+1}(s')) + Q^{\pi^t}(s', \pi^{t+1}(s'))}_{\text{intermediate policy}} - Q^{\pi^t}(s', \pi^t(s')) \right] \\ &\geq \gamma \mathbb{E}_{s' \sim P(s, a)} \left[ Q^{\pi^{t+1}}(s', \pi^{t+1}(s')) - Q^{\pi^t}(s', \pi^{t+1}(s')) \right] \geq \dots, \geq -\gamma^\infty / (1 - \gamma) = 0 \end{aligned}$$



# Monotonic Improvement

$$\pi_{t+1} = \operatorname{argmax}_a Q^{\pi^t}(s, a)$$

---

We want to show that  $Q^{\pi^{t+1}} \geq Q^{\pi^t}$  for all  $s, a$

$$\begin{aligned} Q^{\pi^{t+1}}(s, a) - Q^{\pi^t}(s, a) &= \gamma \mathbb{E}_{s' \sim P(s, a)} \left[ Q^{\pi^{t+1}}(s', \pi^{t+1}(s')) - Q^{\pi^t}(s', \pi^t(s')) \right] \\ &= \gamma \mathbb{E}_{s' \sim P(s, a)} \left[ Q^{\pi^{t+1}}(s', \pi^{t+1}(s')) - Q^{\pi^t}(s', \pi^{t+1}(s')) + \underbrace{Q^{\pi^t}(s', \pi^{t+1}(s')) - Q^{\pi^t}(s', \pi^t(s'))}_{\geq 0} \right] \\ &\geq \gamma \mathbb{E}_{s' \sim P(s, a)} \left[ Q^{\pi^{t+1}}(s', \pi^{t+1}(s')) - Q^{\pi^t}(s', \pi^{t+1}(s')) \right] \geq \dots, \geq -\gamma^\infty / (1 - \gamma) = 0 \end{aligned}$$



# Monotonic Improvement

$$\pi_{t+1} = \operatorname{argmax}_a Q^{\pi^t}(s, a)$$

---

We want to show that  $Q^{\pi^{t+1}} \geq Q^{\pi^t}$  for all  $s, a$

$$\begin{aligned} Q^{\pi^{t+1}}(s, a) - Q^{\pi^t}(s, a) &= \gamma \mathbb{E}_{s' \sim P(s, a)} \left[ Q^{\pi^{t+1}}(s', \pi^{t+1}(s')) - Q^{\pi^t}(s', \pi^t(s')) \right] \\ &= \gamma \mathbb{E}_{s' \sim P(s, a)} \left[ Q^{\pi^{t+1}}(s', \pi^{t+1}(s')) - Q^{\pi^t}(s', \pi^{t+1}(s')) + Q^{\pi^t}(s', \pi^{t+1}(s')) - Q^{\pi^t}(s', \pi^t(s')) \right] \\ &\stackrel{\textcircled{\geq}}{\geq} \gamma \mathbb{E}_{s' \sim P(s, a)} \left[ Q^{\pi^{t+1}}(s', \pi^{t+1}(s')) - Q^{\pi^t}(s', \pi^{t+1}(s')) \right] \geq \dots, \geq -\gamma \textcircled{\infty} (1 - \gamma) = \textcircled{0} \end{aligned}$$



# Properties of Policy Iteration

— — —

- Monotonic improvement:  $Q^{\pi^{t+1}} \geq Q^{\pi^t}$  for all  $s, a$
- Convergence:  $\|V^{\pi^i} - V^*\| \leq \gamma^{i+1} \|V^{\pi^0} - V^*\|$

**Convergence? Prove it yourselves!**



# Properties of Policy Iteration

— — —

- Monotonic improvement:  $Q^{\pi^{t+1}} \geq Q^{\pi^t}$  for all  $s, a$
- Convergence:  $\|V^{\pi^i} - V^*\| \leq \gamma^{i+1} \|V^{\pi^0} - V^*\|$

Complexity  $O(S^3 + S^2A)$



# Properties of Policy Iteration

---

- Monotonic improvement:  $Q^{\pi^{t+1}} \geq Q^{\pi^t}$  for all  $s, a$
- Convergence:  $\|V^{\pi^i} - V^*\| \leq \gamma^{i+1} \|V^{\pi^0} - V^*\|$

**Is there a max number of iterations of policy iteration?**

$|A|^{|S|}$  since that is the maximum number of policies, and as the policy improvement step is monotonically improving, each policy can only appear in one round of policy iteration unless it is an optimal policy



# Properties of Policy Iteration

— — —

- Monotonic improvement:  $Q^{\pi^{t+1}} \geq Q^{\pi^t}$  for all  $s, a$
- Convergence:  $\|V^{\pi^i} - V^*\| \leq \gamma^{i+1} \|V^{\pi^0} - V^*\|$

**When do we stop?**

if the policy does not change anymore for any state



# We Did Dynamic Programming!

---

Dynamic Programming is a method for solving complex problems by breaking them down into subproblems:

- Solve the subproblems
- Combine solutions to subproblems





# We Did Dynamic Programming!

---

Dynamic Programming can be applied if we have:

- *Optimal substructure*: Optimality exists and the optimal solution can be decomposed into subproblems
- *Overlapping subproblems*: Subproblems recur many times and the solutions can be cached and reused

**MDPs satisfy both properties: thanks Bellman equation!**

# We Did Dynamic Programming!

---

We applied dynamic programming for **planning** as we assumed to know the MDP transition probabilities

Problem	Bellman Equation	Algorithm
Prediction	Bellman Expectation Equation	Iterative Policy Evaluation
Control	Bellman Expectation Equation + Greedy Policy Improvement	Policy Iteration
Control	Bellman Optimality Equation	Value Iteration

Credits: David Silver



SAPIENZA  
UNIVERSITÀ DI ROMA

# Primal Linear Program

---

*As an alternative to VI and PI*

Consider the Bellman optimality equation

$$V(s) = \max_a \{r_t + \gamma \mathbb{E}_{s' \sim p(\cdot | s, \pi(s))} [V(s')]\}$$

and write it as a linear program:

$$\min V(s)$$

such that  $V(s) \geq r_t + \gamma \mathbb{E}_{s' \sim p(\cdot | s, \pi(s))} [V(s')] \text{ for all } s, a$



# Primal Linear Program

---

$$\min V(s)$$

such that  $V(s) \geq r_t + \gamma \mathbb{E}_{s' \sim p(\cdot | s, \pi(s))} [V(s')] \text{ for all } s, a$

Using a LP solver we can get a solution which is  $V^*$

**(not used a lot in practice)**



# Primal Linear Program

---

$$\min V(s)$$

such that  $V(s) \geq F(V)$  for all  $s, a$

Any feasible solution must satisfy  $V \geq F(V) \geq F(F(V)) \geq \dots \geq F^\infty V \geq V^*$



# Dual Linear Program

---

There is also a dual linear program, that finds the solution  
directly in policy space

